Artificial Intelligence Project: Hierarchy Deep Q-Learning

Zhihao Jiang, Yiheng Lin

Introduction

Meta Controller:
(Higher Hierarchy)
Interact with External Environment (extrinsic reward), set goals for Controller;

Controller:
(Lower Hierarchy)
Try to achieve goals, receive intrinsic reward from Meta Controller.

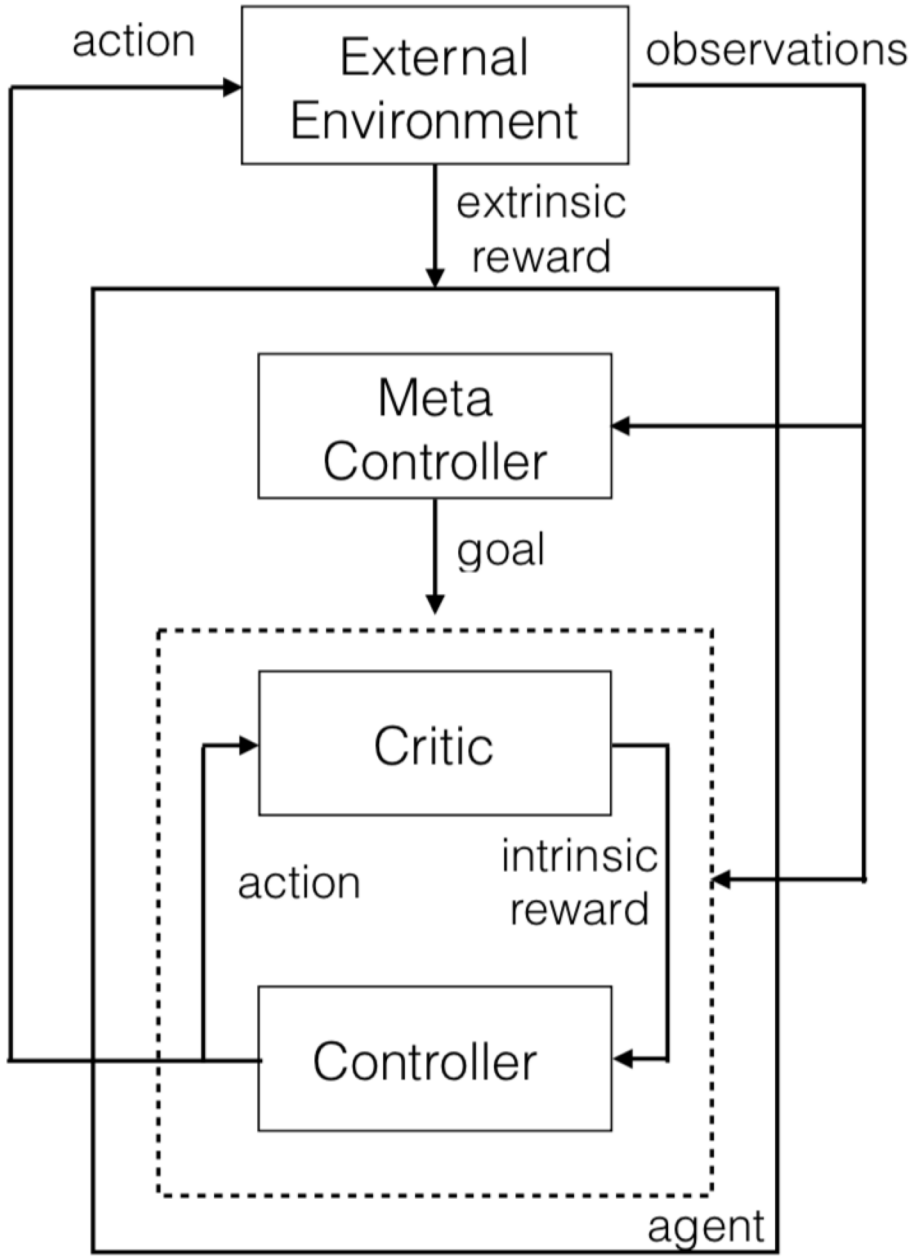

Problem Setting

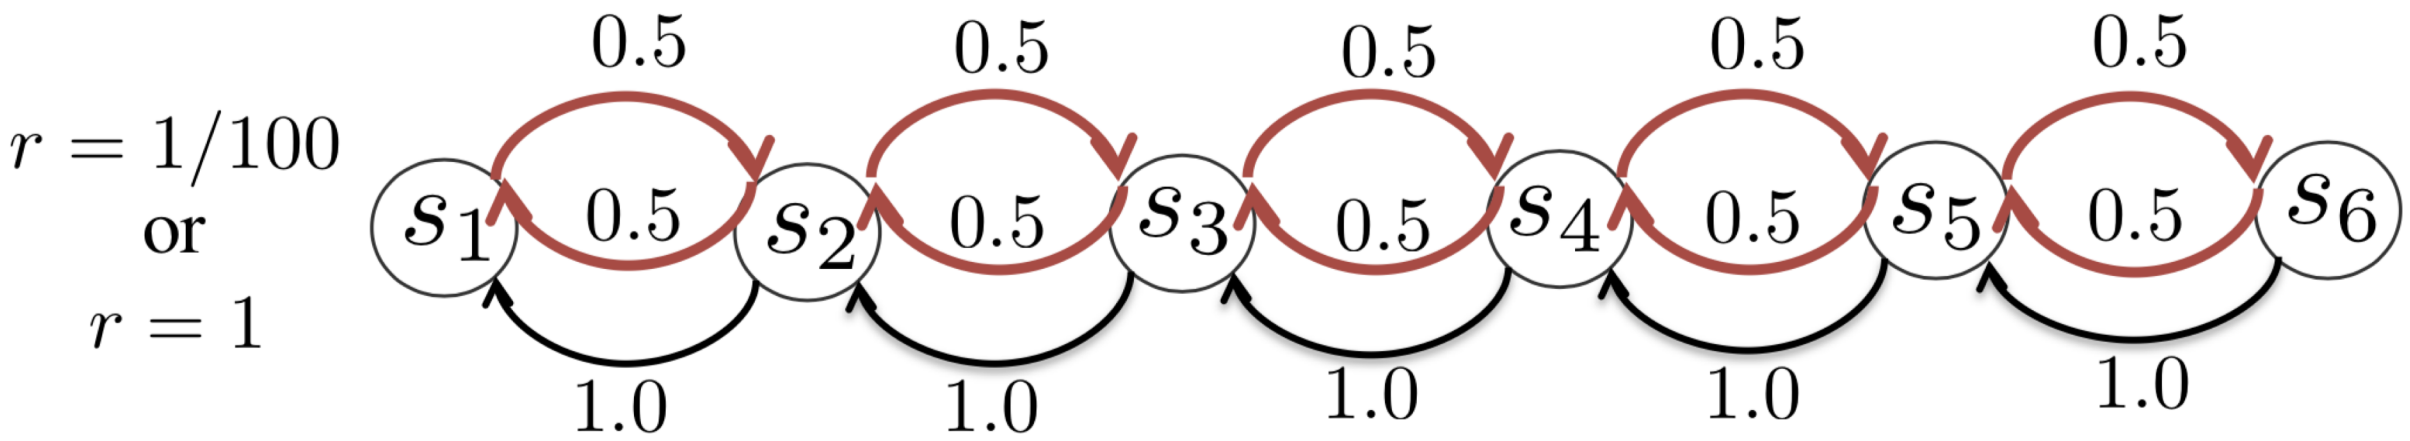

Reward: if s6 is visited, reward = 1; else, reward = 0.01

Actions: 1: move to left with probability 1;

2: move to right with probability 0.5; otherwise, move to left;

Difficulties

$r = 1/100$
or
$r = 1$

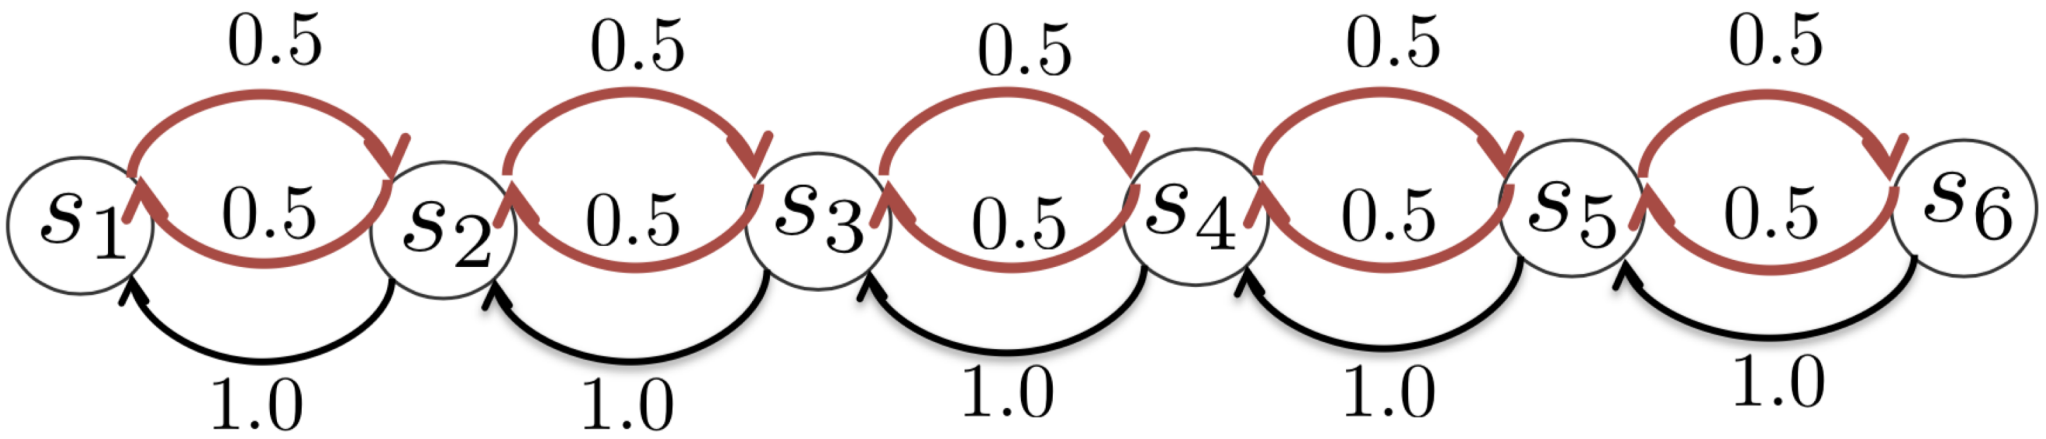

Hidden State: What if we have an additional state to indicate whether s6 is visited?

Search Efficiency: Can epsilon greedy agent keep taking action 2 for enough times?

Reimplementation

3: **for** $i = 1, num\_episodes$ **do**
4: Initialize game and get start state description s
5: $g \leftarrow \text{EPSGREEDY}(s, \mathcal{G}, \epsilon_2, Q_2)$
6: **while** s is **not** terminal **do**
7: $F \leftarrow 0$
8: $s_0 \leftarrow s$
9: **while not** (s is terminal **or** goal g reached) **do**
10: $a \leftarrow \text{EPSGREEDY}(\{s, g\}, \mathcal{A}, \epsilon_{1,g}, Q_1)$
11: Execute a and obtain next state s' and extrinsic reward f from environment
12: Obtain intrinsic reward $r(s, a, s')$ from internal critic
13: Store transition $(\{s, g\}, a, r, \{s', g\})$ in $\mathcal{D}_1$
14: $\text{UPDATEPARAMS}(\mathcal{L}_1(\theta_{1,i}), \mathcal{D}_1)$
15: $\text{UPDATEPARAMS}(\mathcal{L}_2(\theta_{2,i}), \mathcal{D}_2)$
16: $F \leftarrow F + f$
17: $s \leftarrow s'$
18: **end while**
19: Store transition (s_0, g, F, s') in $\mathcal{D}_2$
20: **if** s is **not** terminal **then**
21: $g \leftarrow \text{EPSGREEDY}(s, \mathcal{G}, \epsilon_2, Q_2)$
22: **end if**
23: **end while**

Meta Controller chooses a goal

Controller chooses an action

Controller's reward

Reimplementation

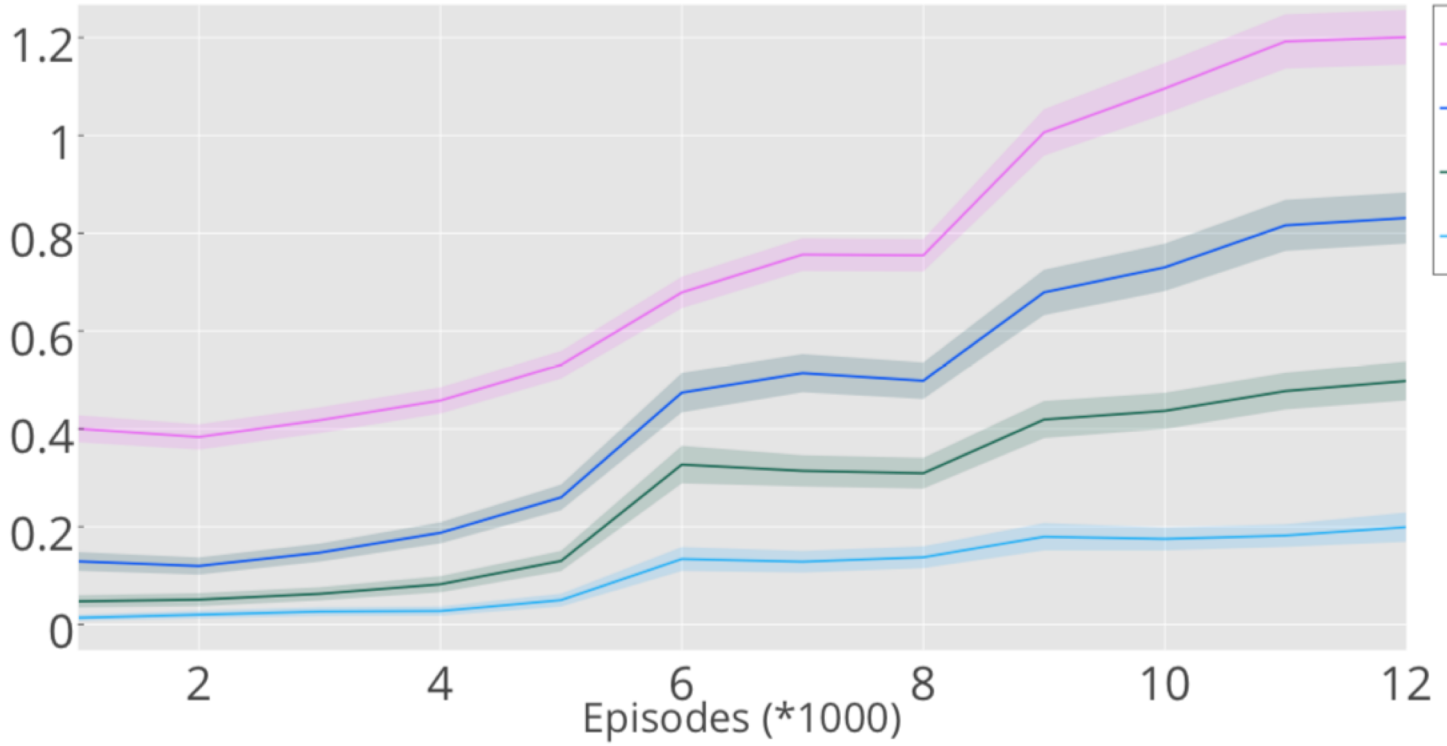

Authors' Result:

Ref: Hierarchical Deep Reinforcement Learning: Integrating Temporal Abstraction and Intrinsic Motivation

Reimplementation

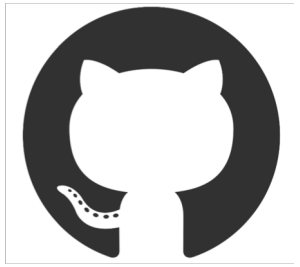

Our result:

epoch	s1	s2	s3	s4	s5	s6
1	1	1.596	0.903	0.421	0.172	0.058
2	1	1.571	0.859	0.435	0.21	0.063
3	1	1.654	0.928	0.415	0.211	0.07
4	1	1.661	0.978	0.484	0.237	0.07
5	1	1.592	0.9	0.499	0.298	0.107
6	1	1.635	1.003	0.601	0.36	0.127
7	1	1.576	0.942	0.574	0.306	0.098

Ref: https://github.com/EthanMacdonald/h-DQN

Reimplementation

- Interesting Phenomena:

When goal is sited to s2, the controller often takes action 2…

Then it can visit s6…

Meta-Controller receives a big reward…

Meta-Controller tends to set s2 as the goal…

But, is this stable?

Maybe the experience of Controller is too volatile to train Meta Controller?

Quality input guarantees quality output…